

# ***ItsRunTym***

## **Operating System**

### **50+ interview questions/answers**

#### **Operating System Basics:**

1. What is an operating system?

An operating system (OS) is a software that acts as an intermediary between computer hardware and user applications. It manages computer hardware resources, provides common services for computer programs, and enables users to interact with the computer.

2. Explain the main functions of an operating system.

The main functions of an operating system include:

- Process management
- Memory management
- File system management
- Device management
- User interface
- Security and access control

3. Describe the difference between a process and a thread.

A process is an instance of a program that is being executed. It has its own memory space, resources, and execution context. A thread is a subset of a process and represents a single flow of execution within that process. Multiple threads within the same process share the same memory space and resources.

4. What are the differences between multiprogramming, multitasking, and multiprocessing?

**Multiprogramming:** In multiprogramming, multiple programs reside in memory simultaneously, and the CPU switches between them to execute, improving CPU utilization.

**Multitasking:** Multitasking refers to the ability of an operating system to execute multiple tasks (programs or processes) concurrently, allowing users to run multiple applications simultaneously.

**Multiprocessing:** Multiprocessing involves the use of multiple CPUs or cores in a computer system to execute multiple processes concurrently, improving overall system performance.

5. Explain the concept of a context switch.

A context switch is the process of saving the state of a running process or thread so that it can be restored and resumed later, while another process or thread is given the CPU. It involves saving the current execution context, including program counter, registers, and stack pointer, and loading the context of the next process or thread to be executed.

6. What are the differences between a monolithic kernel and a microkernel?

**Monolithic Kernel:** In a monolithic kernel, all operating system services run in a single address space, and system calls directly invoke kernel functions. Examples include Linux and Unix.

**Microkernel:** In a microkernel architecture, only essential functions, such as memory management and inter-process communication, run in kernel space. Other services, like device drivers and file systems, run as user-space processes. Examples include QNX and Minix.

7. Describe the process of process creation and termination.

Process creation involves allocating resources, such as memory and CPU time, to a new process. This includes initializing process control blocks, assigning process IDs, and setting up execution contexts. Process termination involves releasing allocated resources, closing files, and deallocating memory associated with the process.

8. What is the difference between preemptive and non-preemptive scheduling?

Preemptive scheduling allows a higher priority task to interrupt a lower priority task currently executing.

In non-preemptive scheduling, a task holds the CPU until it voluntarily relinquishes it or completes its execution.

9. What are system calls, and how are they different from normal function calls?

System calls are interface functions provided by the operating system that allow user-level processes to request services from the operating system kernel. They provide access to operating system services such as file operations, process management, and inter-process communication.

System calls transition the processor from user mode to kernel mode to execute privileged instructions, whereas normal function calls execute within the user space.

10. Explain the concept of kernel mode and user mode.

**Kernel Mode:** Also known as privileged mode or supervisor mode, the kernel mode is a privileged execution mode where the CPU can execute privileged instructions and access protected resources such as hardware devices and memory management. Operating system code runs in kernel mode.

**User Mode:** In user mode, user-level applications execute. They have limited access to system resources and cannot directly

execute privileged instructions or access hardware devices. Most user applications run in user mode.

## **Process Management:**

### 1. What is a microprocessor?

Computer's Central Processing Unit (CPU) built on a **single Integrated Circuit (IC)** is called a **microprocessor**.

A digital computer with one microprocessor which acts as a CPU is called microcomputer.

It is a programmable, multipurpose, clock -driven, register-based electronic device that reads binary instructions from a storage device called memory, accepts binary data as input and processes data according to those instructions and provides results as output.

### 2. What is process scheduling.

Process scheduling is the activity of the process manager that handles the removal of the running process from the CPU and the selection of another process based on a particular strategy.

Process scheduling is an essential part of a Multiprogramming operating system. Such operating systems allow more than one process to be loaded into the executable memory at a time and the loaded process shares the CPU using time multiplexing

### 3. What are the different scheduling algorithms used in operating systems?

- First-Come, First-Served (FCFS) Scheduling
- Shortest-Job-Next (SJN) Scheduling
- Priority Scheduling
- Shortest Remaining Time
- Round Robin(RR) Scheduling
- Multiple-Level Queues Scheduling

4. Explain the differences between preemptive and non-preemptive scheduling.

### **Preemptive Scheduling**

Preemptive scheduling is used when a process switches from the running state to the ready state or from the waiting state to the ready state. The resources (mainly CPU cycles) are allocated to the process for a limited amount of time and then taken away, and the process is again placed back in the ready queue if that process still has CPU burst time remaining. That process stays in the ready queue till it gets its next chance to execute.

### **Non-Preemptive Scheduling**

Non-preemptive Scheduling is used when a process terminates, or a process switches from running to the waiting state. In this scheduling, once the resources (CPU cycles) are allocated to a process, the process holds the CPU till it gets terminated or reaches a waiting state. In the case of non-preemptive scheduling does not interrupt a process running CPU in the middle of the execution. Instead, it waits till the process completes its CPU burst time, and then it can allocate the CPU to another process.

5. What is a context switch, and how does it affect the performance of a system?

A context switch is the process of saving the state of a running process or thread so that it can be restored and resumed later, while another process or thread is given the CPU. It affects system performance by consuming CPU cycles and memory resources.

6. Describe the process of process synchronization using semaphores.

Process synchronization using semaphores involves using synchronization primitives called semaphores to coordinate access to shared resources among multiple processes or threads. Semaphores can be used to control access to critical sections of code and prevent race conditions.

7. Explain the dining philosophers' problem and how it can be solved.

The dining philosophers' problem is a classic synchronization problem where a group of philosophers sit at a table with a bowl of rice. Between each philosopher, there is a single chopstick. The problem arises when each philosopher alternates between thinking and eating, requiring two chopsticks to eat. If each philosopher tries to pick up both chopsticks simultaneously, they may end up in a deadlock. This problem can be solved using various synchronization techniques like mutexes or semaphores to ensure that philosophers can only eat when both chopsticks are available.

8. What is a critical section, and how is it protected in concurrent programming?

A critical section is a segment of code in which shared resources are accessed and manipulated. In concurrent programming, a critical section needs to be protected to prevent race conditions and ensure data integrity. This is typically done using synchronization mechanisms such as locks, semaphores, or mutexes to allow only one process or thread to access the critical section at a time.

9. Explain the reader-writer problem and how it can be solved.

The reader-writer problem involves multiple processes or threads accessing a shared resource, where some only read the resource (readers) while others both read and write to it (writers). The problem is to ensure that multiple readers can access the resource simultaneously but exclusive access is granted to writers to maintain data consistency. Solutions often involve using semaphore-based synchronization to prioritize access based on the type of operation (read or write).

10. Describe the process of process communication using inter-process communication (IPC).

Process communication using Inter-Process Communication (IPC) involves mechanisms for data exchange and synchronization between processes running on a computer system. This can include methods such as pipes, message queues, shared memory, and sockets.

11. What are the different IPC mechanisms available in operating systems?

- Pipes
- Message Queues
- Shared Memory
- Semaphores
- Sockets
- Remote Procedure Calls (RPC)
- Signals

## **Memory Management:**

1. What is virtual memory, and how does it work?

Virtual memory is a memory management technique that provides an illusion of a larger memory space than physically available in the system. It allows programs to use more memory than is actually installed on the computer by temporarily transferring data from RAM to disk.

Virtual memory works by dividing the physical memory into fixed-sized blocks called pages. When a process requests memory, the operating system allocates space in virtual memory, which may not necessarily be in RAM at that moment. Pages are then swapped between RAM and disk as needed.

2. Explain the concept of paging and its advantages.

Paging is a memory management scheme that divides the virtual memory and physical memory into fixed-size blocks called

pages. Similarly, the process memory is divided into fixed-size blocks called page frames. Paging allows non-contiguous allocation of memory, efficient use of physical memory, and simplifies memory management compared to contiguous memory allocation.

3. What is a page fault, and how is it handled by the operating system?

A page fault occurs when a program accesses a page of memory that is not currently in RAM (i.e., it's stored on disk). When a page fault occurs, the operating system handles it by:

- Identifying the page that needs to be brought into RAM.
- Selecting a page frame in RAM to hold the new page.
- Loading the requested page from disk into the selected page frame.
- Updating the page table to reflect the new mapping.
- Resuming the interrupted process.

4. Describe the process of memory allocation and deallocation.

Memory allocation involves allocating memory space to processes or data structures dynamically. This typically involves requesting memory from the operating system using functions like `malloc()` or `new` in languages like C and C++.

Memory deallocation involves releasing allocated memory when it is no longer needed, preventing memory leaks and fragmentation. This is done using functions like `free()` or `delete`.

5. Explain the concepts of thrashing and working set model.

- Thrashing: Thrashing occurs when the system spends a significant amount of time swapping pages between RAM and disk due to excessive paging activity. It leads to a decrease in system performance as the CPU spends more



time on disk I/O operations rather than executing processes.

- **Working Set Model:** The working set model is a concept used to describe the set of pages that a process actively uses during its execution. It helps in determining the amount of memory a process requires to run efficiently without excessive paging. By keeping the working set in RAM, thrashing can be minimized.

6. Describe the different page replacement algorithms, such as LRU, FIFO, and Optimal.

- **LRU (Least Recently Used):** LRU replaces the page that has not been accessed for the longest period of time.
- **FIFO (First In, First Out):** FIFO replaces the oldest page in memory, based on the order they were brought in.
- **Optimal:** Optimal replacement replaces the page that will not be used for the longest period of time in the future. However, it's not practical as it requires knowledge of future memory accesses.

7. What is the purpose of a page table, and how is it used in virtual memory management?

A page table is a data structure used by the operating system to map virtual addresses to physical addresses in virtual memory management. It stores the mapping between virtual pages and physical page frames, along with additional information such as page status and permissions.

When a program accesses memory, the page table is consulted to translate the virtual address to the corresponding physical address.

8. Explain the concept of demand paging and its advantages.

Demand paging is a technique where pages are only loaded into memory when they are needed, rather than loading the entire program into memory at once. This reduces the initial loading time and conserves memory resources since only the necessary pages are brought into memory as required.

9. What is a segmentation fault, and how is it handled by the operating system?

A segmentation fault, also known as a segfault, occurs when a program attempts to access memory that it does not have permission to access, or memory that is not mapped to it.

When a segmentation fault occurs, the operating system typically terminates the offending process and may generate an error message or core dump for debugging purposes.

10. Describe the process of process swapping.

Process swapping is a technique used by the operating system to move entire processes between main memory (RAM) and secondary storage (usually disk) to free up memory for other processes.

When a process is swapped out, its entire memory image is written to disk. When it needs to be run again, it is swapped back into memory. Swapping allows the system to handle more processes than can fit into physical memory at once, improving overall system performance.

## **File Systems:**

1. What is a file system, and what are its components?

A file system is a method of organizing and storing computer data on a storage device, such as a hard disk drive or solid-state drive. Its components include:

- File: A named collection of related data stored on a storage medium.
  - Directory: A container for files and subdirectories, used for organizing and managing files.
  - Metadata: Information about files and directories, such as file size, permissions, and timestamps.
  - File system operations: Functions for creating, reading, writing, and deleting files.
2. Explain the different types of file systems, such as FAT, NTFS, and ext4.
- FAT (File Allocation Table): A simple file system used primarily on older Windows operating systems. It uses a table to allocate space for files on disk.
  - NTFS (New Technology File System): A modern file system developed by Microsoft for Windows NT and later versions. It offers features like improved security, file compression, and support for large file sizes and volumes.
  - ext4: The fourth extended file system, commonly used in Linux distributions. It provides improvements over its predecessor, ext3, including support for larger file sizes and volumes, faster file system checks, and better performance.
3. Describe the process of file allocation and deallocation.

File allocation involves allocating space on a storage device to store a file's data. This typically involves finding free space on the disk and updating the file system's data structures to record the allocation.

File deallocation involves releasing the allocated space when a file is deleted or its size is reduced. This space is marked as free in the file system's data structures, making it available for future allocation.

4. What is a file control block (FCB) or an inode, and how is it used in file systems?

A file control block (FCB) or an inode (index node) is a data structure used by file systems to store metadata about files. It contains information such as file size, permissions, timestamps, and pointers to data blocks or extents where the file's data is stored on disk. The FCB or inode is used by the operating system to manage and access files efficiently.

5. Explain the concepts of file descriptors and file descriptor tables.

- File Descriptor: A unique identifier assigned by the operating system to an open file. It represents a reference to the file in the operating system's file management system.
- File Descriptor Table: A data structure maintained by the operating system for each process, containing entries for all open files by that process. Each entry in the table corresponds to a file descriptor and contains information about the file, such as its current position and access mode.

6. What is a file allocation table (FAT), and how does it work?

A file allocation table (FAT) is a data structure used by the FAT file system to maintain a map of the clusters on a disk and their allocation status (i.e., whether they are free or allocated to files).

The FAT contains entries for each cluster on the disk, indicating whether the cluster is available or allocated to a file. It is used by the file system to allocate space for files and to locate the clusters that contain a file's data.

7. Describe the differences between sequential, direct, and indexed file allocation methods.

- Sequential Allocation: Files are allocated contiguous blocks of disk space. It is simple but can lead to fragmentation.
- Direct Allocation: Each file has a list of pointers to its data blocks. It reduces fragmentation but has limited scalability.
- Indexed Allocation: Each file has an index block containing pointers to its data blocks. It allows for dynamic file size and efficient random access.

8. Explain the concept of file buffering and its advantages.

File buffering involves temporarily storing data in memory buffers before reading from or writing to a file on disk.

It improves I/O performance by reducing the number of disk accesses and optimizing data transfers between the CPU and storage device. Buffered I/O can significantly speed up file operations, especially for small reads and writes.

9. What is a symbolic link, and how does it work in file systems?

A symbolic link, also known as a symlink or soft link, is a special type of file that points to another file or directory in the file system. Unlike hard links, symbolic links are references to the file's path rather than its inode.

They provide flexibility and convenience by allowing files to be accessed from multiple locations without duplicating data.

10. Describe the process of file permission management in operating systems.

File permission management involves assigning permissions to files and directories to control access by users and groups.

Permissions typically include read, write, and execute permissions for the owner, group, and others. The operating system enforces these permissions when users attempt to access or modify files, ensuring data security and integrity.

Administrators can use commands like `chmod` and `chown` to modify file permissions and ownership.

## **Device Management:**

1. What is a device driver, and what is its role in an operating system?

A device driver is a specialized software component that enables communication between the operating system and hardware devices attached to the computer system, such as printers, disk drives, and network adapters.

Its role is to abstract the hardware-specific details of the device and provide a standardized interface for the operating system to interact with the device.

2. Explain the process of device allocation and deallocation.

Device allocation involves assigning a device to a process or application for exclusive or shared use. This typically involves identifying an available device, checking its status and compatibility with the requesting process, and allocating the necessary resources.

Device deallocation involves releasing the device resources when they are no longer needed, allowing other processes to use the device.

3. What are the different types of device scheduling algorithms used in operating systems?

Device scheduling algorithms determine the order in which device requests are serviced by the operating system. Common device scheduling algorithms include:

- First Come, First Served (FCFS)

- Shortest Seek Time First (SSTF)
- SCAN
- C-SCAN
- LOOK
- C-LOOK

4. Describe the process of device interrupt handling.

Device interrupt handling involves the operating system responding to hardware interrupts generated by devices to signal events such as data transfer completion or error conditions.

When an interrupt occurs, the CPU suspends its current execution, saves the context, and jumps to the interrupt handler routine specified by the interrupt vector table.

The interrupt handler then processes the interrupt, performs any necessary actions, and returns control to the interrupted process.

5. What is a device control block (DCB), and how is it used in device management?

A device control block (DCB) is a data structure maintained by the operating system to manage information about a specific device, such as its status, configuration, and current state.

DCBs are used by the device driver to communicate with the device and by the operating system to coordinate device operations and handle device requests from processes.

6. Explain the concept of spooling and its benefits.

Spooling (Simultaneous Peripheral Operation On-Line) is a technique used to improve the performance and efficiency of I/O operations by storing data temporarily in a buffer or queue before processing.

It allows multiple processes to send their output to a common spooling area, which is then processed sequentially by a single device, such as a printer.

Spooling helps to smooth out variations in I/O rates and prevents processes from blocking while waiting for slow devices, improving overall system throughput.

7. What is a device register, and how does it relate to device management?

A device register is a hardware component within a device that stores control and status information used for device management. It contains information such as device configuration settings, operational mode, and error flags.

Device drivers interact with device registers to configure the device, initiate data transfers, and monitor device status during I/O operations.

8. Describe the differences between polling and interrupt-driven I/O.

- Polling: In polling, the CPU continuously checks the status of the device at regular intervals to determine if it is ready for data transfer. Polling is simple but can lead to high CPU overhead and inefficiency.
- Interrupt-driven I/O: In interrupt-driven I/O, the device signals the CPU with an interrupt when it is ready for data transfer or when an error occurs. The CPU responds to the interrupt by suspending its current execution and servicing the device request. Interrupt-driven I/O is more efficient as it eliminates the need for constant CPU polling.

9. What is a device queue, and how is it used in device management?



A device queue is a data structure used by the operating system to manage pending device requests from processes. It maintains a queue of requests for each device, ordering them based on priority or arrival time.

The device driver services requests from the device queue in the order specified by the scheduling algorithm, ensuring fair access to the device and efficient resource utilization.

10. Explain the concept of device management.

Device management involves the coordination and control of hardware devices attached to a computer system, including initialization, configuration, allocation, and deallocation of devices.

It encompasses tasks such as device detection, driver installation, device status monitoring, error handling, and resource arbitration to ensure efficient and reliable operation of the system.

Device management is essential for maintaining system stability, performance, and compatibility with a wide range of hardware configurations.